

Loading all the Libraries

In []:

```
from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

In []:

```
!pip install transformers==3.5.1
```

Collecting transformers==3.5.1

Downloading

<https://files.pythonhosted.org/packages/3a/83/e74092e7f24a08d751aa59b37a9fc572b2e4af3918cb66f7766c31b4/transformers-3.5.1-py3-none-any.whl> (1.3MB)

|██| 1.3MB 7.4MB/s

Requirement already satisfied: dataclasses; python_version < "3.7" in /usr/local/lib/python3.6/dist-packages (from transformers==3.5.1) (0.8)

Requirement already satisfied: regex!=2019.12.17 in /usr/local/lib/python3.6/dist-packages (from transformers==3.5.1) (2019.12.20)

Requirement already satisfied: tqdm>=4.27 in /usr/local/lib/python3.6/dist-packages (from transformers==3.5.1) (4.41.1)

Requirement already satisfied: protobuf in /usr/local/lib/python3.6/dist-packages (from transformers==3.5.1) (3.12.4)

Collecting tokenizers==0.9.3

Downloading

https://files.pythonhosted.org/packages/4c/34/b39eb9994bc3c999270b69c9eea40ecc6f0e97991dba28282b9fc4ee/tokenizers-0.9.3-cp36-cp36m-manylinux1_x86_64.whl (2.9MB)

|██| 2.9MB 24.5MB/s

Requirement already satisfied: packaging in /usr/local/lib/python3.6/dist-packages (from transformers==3.5.1) (20.4)

Requirement already satisfied: numpy in /usr/local/lib/python3.6/dist-packages (from transformers==3.5.1) (1.18.5)

Requirement already satisfied: filelock in /usr/local/lib/python3.6/dist-packages (from transformers==3.5.1) (3.0.12)

Requirement already satisfied: requests in /usr/local/lib/python3.6/dist-packages (from transformers==3.5.1) (2.23.0)

Collecting sentencepiece==0.1.91

Downloading

https://files.pythonhosted.org/packages/d4/a4/d0a884c4300004a78cca907a6ff9a5e9fe4f090f5d95ab341c53cc58/sentencepiece-0.1.91-cp36-cp36m-manylinux1_x86_64.whl (1.1MB)

|██| 1.1MB 42.3MB/s

Collecting sacremoses

Downloading

<https://files.pythonhosted.org/packages/7d/34/09d19aff26edcc8eb2a01bed8e98f13a1537005d31e95233fd482d10/sacremoses-0.0.43.tar.gz> (883kB)

|██| 890kB 45.7MB/s

Requirement already satisfied: setuptools in /usr/local/lib/python3.6/dist-packages (from protobuf->transformers==3.5.1) (50.3.2)

Requirement already satisfied: six>=1.9 in /usr/local/lib/python3.6/dist-packages (from protobuf->transformers==3.5.1) (1.15.0)

Requirement already satisfied: pyparsing>=2.0.2 in /usr/local/lib/python3.6/dist-packages (from packaging->transformers==3.5.1) (2.4.7)

Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.6/dist-packages (from requests->transformers==3.5.1) (2020.11.8)

Requirement already satisfied: idna<3,>=2.5 in /usr/local/lib/python3.6/dist-packages (from requests->transformers==3.5.1) (2.10)

Requirement already satisfied: urllib3!=1.25.0,!<1.25.1,<1.26,>=1.21.1 in /usr/local/lib/python3.6/dist-packages (from requests->transformers==3.5.1) (1.24.3)

Requirement already satisfied: chardet<4,>=3.0.2 in /usr/local/lib/python3.6/dist-packages (from requests->transformers==3.5.1) (3.0.4)

Requirement already satisfied: click in /usr/local/lib/python3.6/dist-packages (from sacremoses->transformers==3.5.1) (7.1.2)

Requirement already satisfied: joblib in /usr/local/lib/python3.6/dist-packages (from sacremoses->transformers==3.5.1) (0.17.0)

Building wheels for collected packages: sacremoses

Building wheel for sacremoses (setup.py) ... done

Created wheel for sacremoses: filename=sacremoses-0.0.43-cp36-none-any.whl size=893257

```
sha256=fe738906e00c5c2ace1e047302f8ac56519bbb0ec6381086b50e3026b4e80f9c
```

```
Stored in directory:
/root/.cache/pip/wheels/29/3c/fd/7ce5c3f0666dab31a50123635e6fb5e19ceb42ce38d4e58f45
Successfully built sacremoses
Installing collected packages: tokenizers, sentencepiece, sacremoses, transformers
Successfully installed sacremoses-0.0.43 sentencepiece-0.1.91 tokenizers-0.9.3 transformers-3.5.1
```

```
In [ ]:
```

```
import pandas as pd
import numpy as np
import re
import math
import random
import tqdm
import pickle
import random
import time
from datetime import datetime
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
import matplotlib.pyplot as plt

#Torch
import torch
import torch.nn as nn
from torch.utils import data
from torch.utils.data import DataLoader
from torchsummary import summary
from torch.utils.data import random_split
from torch.utils.data import TensorDataset
from torch.utils.data import RandomSampler
from torch.utils.data import SequentialSampler
from torch.optim.lr_scheduler import CosineAnnealingLR
from torch.optim.lr_scheduler import ReduceLROnPlateau

#Transformers
from transformers import get_linear_schedule_with_warmup
from transformers import BertForSequenceClassification, AdamW, BertConfig
from transformers import BertTokenizer

#Data
from sklearn.datasets import fetch_20newsgroups
```

Task 1: Classification using BERT or BERT variations

Preparing the Data

```
In [ ]:
```

```
#Train and test dataset
train = fetch_20newsgroups(subset='train', remove=('headers', 'footers', 'quotes'))
test = fetch_20newsgroups(subset='test', remove=('headers', 'footers', 'quotes'))
```

Downloading 20news dataset. This may take a few minutes.
Downloading dataset from <https://ndownloader.figshare.com/files/5975967> (14 MB)

```
In [ ]:
```

```
params = dict(
    seed_val = 1234,
    batch_size = 16,
    path = r'/content/drive/MyDrive')
```

```
In [ ]:
```

```
#Train
df_train = pd.DataFrame()
df_train['data'] = train.data
df_train['labels'] = train.target
```

```
#Test
df_test = pd.DataFrame()
df_test['data'] = test.data
df_test['labels'] = test.target
#df_train.shape, df_test.shape
```

In []:

```
#Train
sentences_train = df_train.data.values
labels_train = df_train.labels.values
#Test
sentences_test = df_test.data.values
labels_test = df_test.labels.values
```

In []:

```
tokenizer = BertTokenizer.from_pretrained('bert-base-uncased', do_lower_case=True)
```

In []:

```
# Print the original sentence.
print(' Original: ', sentences_train[0])
# Print the sentence split into tokens.
print('Tokenized: ', tokenizer.tokenize(sentences_train[0]))
# Print the sentence mapped to token ids.
print('Token IDs: ', tokenizer.convert_tokens_to_ids(tokenizer.tokenize(sentences_train[0])))
```

Original: I was wondering if anyone out there could enlighten me on this car I saw the other day. It was a 2-door sports car, looked to be from the late 60s/early 70s. It was called a Bricklin. The doors were really small. In addition, the front bumper was separate from the rest of the body. This is all I know. If anyone can tell me a model name, engine specs, years of production, where this car is made, history, or whatever info you have on this funky looking car, please e-mail.

Tokenized: ['i', 'was', 'wondering', 'if', 'anyone', 'out', 'there', 'could', 'en', '##light', '##en', 'me', 'on', 'this', 'car', 'i', 'saw', 'the', 'other', 'day', '.', 'it', 'was', 'a', '2', '-', 'door', 'sports', 'car', ',', 'looked', 'to', 'be', 'from', 'the', 'late', '60s', '/', 'early', '70s', '.', 'it', 'was', 'called', 'a', 'brick', '##lin', '.', 'the', 'doors', 'were', 'really', 'small', '.', 'in', 'addition', ',', 'the', 'front', 'bumper', 'was', 'separate', 'from', 'the', 'rest', 'of', 'the', 'body', '.', 'this', 'is', 'all', 'i', 'know', '.', 'if', 'anyone', 'can', 'tell', '##me', 'a', 'model', 'name', ',', 'engine', 'spec', '##s', ',', 'years', 'of', 'production', ',', 'where', 'this', 'car', 'is', 'made', ',', 'history', ',', 'or', 'whatever', 'info', 'you', 'have', 'on', 'this', 'funky', 'looking', 'car', ',', 'please', 'e', '-', 'mail', '.']

Token IDs: [1045, 2001, 6603, 2065, 3087, 2041, 2045, 2071, 4372, 7138, 2368, 2033, 2006, 2023, 2482, 1045, 2387, 1996, 2060, 2154, 1012, 2009, 2001, 1037, 1016, 1011, 2341, 2998, 2482, 1010, 2246, 2000, 2022, 2013, 1996, 2397, 20341, 1013, 2220, 17549, 1012, 2009, 2001, 2170, 1037, 5318, 4115, 1012, 1996, 4303, 2020, 2428, 2235, 1012, 1999, 2804, 1010, 1996, 2392, 21519, 2001, 3584, 2013, 1996, 2717, 1997, 1996, 2303, 1012, 2023, 2003, 2035, 1045, 2113, 1012, 2065, 3087, 2064, 2425, 4168, 1037, 2944, 2171, 1010, 3194, 28699, 2015, 1010, 2086, 1997, 2537, 1010, 2073, 2023, 2482, 2003, 2081, 1010, 2381, 1010, 2030, 3649, 18558, 2017, 2031, 2006, 2023, 24151, 2559, 2482, 1010, 3531, 1041, 1011, 5653, 1012]

All the necessary Functions

In []:

```
#Tokenizer and preprocessing
def tokenizer_s(sentences, max_len):
    input_ids = []
    attention_masks = []
    for sent in sentences:
        encoded_dict = tokenizer.encode_plus(
            sent,
            add_special_tokens = True,
            max_length = max_len,
            pad_to_max_length = True,
            return_attention_mask = True,
```

```

        return_attention_mask = True,
        return_tensors = 'pt',
    )

    # Add the encoded sentence to the list.
    input_ids.append(encoded_dict['input_ids'])

    # And its attention mask (simply differentiates padding from non-padding).
    attention_masks.append(encoded_dict['attention_mask'])

# Convert the lists into tensors.
input_ids = torch.cat(input_ids, dim=0)
attention_masks = torch.cat(attention_masks, dim=0)

return input_ids, attention_masks

```

In []:

```

#data loaders function

def loaders(sentences_train, labels_train, sentences_test, labels_test, max_len, batch_size):
    input_ids_train, attention_masks_train, input_ids_test, attention_masks_test, labels_train, labels_test = sentences_t(sentences_train, labels_train, sentences_test, labels_test, max_len)

    train_dataset = TensorDataset(input_ids_train, attention_masks_train, labels_train)
    test_dataset = TensorDataset(input_ids_test, attention_masks_test, labels_test)

    train_dataloader = DataLoader(
        train_dataset,
        sampler = RandomSampler(train_dataset),
        batch_size = batch_size
    )

    validation_dataloader = DataLoader(
        test_dataset,
        sampler = SequentialSampler(test_dataset),
        batch_size = batch_size
    )

    return train_dataloader, validation_dataloader

```

In []:

```

#Sentences to Tokens conversion
def sentences_t(sentences_train, labels_train, sentences_test, labels_test, max_len):
    input_ids_train, attention_masks_train = tokenizer_s(sentences_train, max_len)
    labels_train = torch.tensor(labels_train)

    input_ids_test, attention_masks_test = tokenizer_s(sentences_test, max_len)
    labels_test = torch.tensor(labels_test)

    # Print sentence 0, now as a list of IDs.
    print('Original: ', sentences_train[0])
    print('Token IDs:', input_ids_train[0].shape)

    return input_ids_train, attention_masks_train, input_ids_test, attention_masks_test, labels_train, labels_test

```

In []:

```

#Train loop
def train_loop(model, train_dataloader, validation_dataloader, optimizer, scheduler, epochs, seed, device):

    random.seed(seed)
    np.random.seed(seed)
    torch.manual_seed(seed)
    torch.cuda.manual_seed_all(seed)

    train_losses= np.zeros(epochs)
    valid_losses= np.zeros(epochs)

    for epoch in range(epochs):

        t0= datetime.now()
        train_loss=[]

```

```

train_loss=[]

model.train()
for batch in train_dataloader:

    b_input_ids = batch[0].to(device)
    b_input_mask = batch[1].to(device)
    b_labels = batch[2].to(device)
    # forward pass

    loss, logits = model(b_input_ids,
                          token_type_ids=None,
                          attention_mask=b_input_mask,
                          labels=b_labels)

    # set gradients to zero
    optimizer.zero_grad()

    # backward pass
    loss.backward()
    torch.nn.utils.clip_grad_norm_(model.parameters(), 1.0)
    optimizer.step()
    scheduler.step()
    train_loss.append(loss.item())

train_loss=np.mean(train_loss)

valid_loss=[]
model.eval()
with torch.no_grad():
    for batch in validation_dataloader:

        # forward pass
        b_input_ids = batch[0].to(device)
        b_input_mask = batch[1].to(device)
        b_labels = batch[2].to(device)

        loss, logits = model(b_input_ids,
                              token_type_ids=None,
                              attention_mask=b_input_mask,
                              labels=b_labels)

        valid_loss.append(loss.item())

    valid_loss=np.mean(valid_loss)

# save Losses
train_losses[epoch]= train_loss
valid_losses[epoch]= valid_loss
dt= datetime.now()-t0
print(f'Epoch {epoch+1}/{epochs}, Train Loss: {train_loss:.4f}    Valid Loss: {valid_loss:.4f},
Duration: {dt}')

return train_losses, valid_losses

```

In []:

```

# Accuracy- write a function to get accuracy
# use this function to get accuracy and print accuracy
def get_accuracy(data_iter, model, device):
    model.eval()
    with torch.no_grad():
        correct =0
        total =0

    for batch in data_iter:

        b_input_ids = batch[0].to(device)
        b_input_mask = batch[1].to(device)
        b_labels = batch[2].to(device)
        # forward pass

        loss, logits = model(b_input_ids,
                              token_type_ids=None,
                              attention_mask=b_input_mask,
                              labels=b_labels)

```

```

_,indices = torch.max(logits,dim=1)
correct+= (b_labels==indices).sum().item()
total += b_labels.shape[0]

acc= correct/total

return acc

```

In []:

```

# Write a function to get predictions
def get_predictions(data_iter, model):
    model.eval()
    with torch.no_grad():
        predictions= np.array([])
        y_test= np.array([])

        for batch in data_iter:

            b_input_ids = batch[0].to(device)
            b_input_mask = batch[1].to(device)
            b_labels = batch[2].to(device)
            # forward pass

            loss, logits = model(b_input_ids,
                                token_type_ids=None,
                                attention_mask=b_input_mask,
                                labels=b_labels)

            _,indices = torch.max(logits,dim=1)
            predictions=np.concatenate((predictions,indices.cpu().numpy()))
            y_test = np.concatenate((y_test,b_labels.cpu().numpy()))

    return y_test, predictions

```

In []:

```

def model_parameters(model):
    epochs = 2
    no_decay = ['bias', 'LayerNorm.weight']
    optimizer_grouped_parameters = [
        {'params': [p for n, p in model.named_parameters()
                    if not any(nd in n for nd in no_decay)],
         'weight_decay': 0.5},

        {'params': [p for n, p in model.named_parameters()
                    if any(nd in n for nd in no_decay)],
         'weight_decay': 0.0}
    ]

    optimizer = AdamW(optimizer_grouped_parameters,
                      lr = 5e-5,
                      eps = 1e-8
                      )

    scheduler = get_linear_schedule_with_warmup(optimizer,
                                                num_warmup_steps = 0,
                                                num_training_steps = len(train_dataloader) * epochs)

    return epochs, optimizer, scheduler

```

In []:

```
device = torch.device('cuda:0' if torch.cuda.is_available() else 'cpu')
```

EXPERIMENT 1

BERT Model with 512 tokens

In []:

```
""" [ ]:
```

```
params['max_len'] = 0
for i,sent in enumerate(sentences_train):
    input_ids = tokenizer.encode(sent, add_special_tokens=True)
    params['max_len'] = max(params['max_len'], len(input_ids))
print('Max sentence length: ', params['max_len'])
params['max_len'] = 512
```

Token indices sequence length is longer than the specified maximum sequence length for this model (604 > 512). Running this sequence through the model will result in indexing errors

Max sentence length: 52886

```
In [ ]:
```

```
#Train Data Loader
train_dataloader, validation_dataloader = loaders(sentences_train, labels_train, sentences_test, labels_test, params['max_len'], params['batch_size'])
```

Truncation was not explicitly activated but `max_length` is provided a specific value, please use `truncation=True` to explicitly truncate examples to max length. Defaulting to 'longest_first' truncation strategy. If you encode pairs of sequences (GLUE-style) with the tokenizer you can select this strategy more precisely by providing a specific strategy to `truncation`.

/usr/local/lib/python3.6/dist-packages/transformers/tokenization_utils_base.py:2022:

FutureWarning: The `pad_to_max_length` argument is deprecated and will be removed in a future version, use `padding=True` or `padding='longest'` to pad to the longest sequence in the batch, or use `padding='max_length'` to pad to a max length. In this case, you can give a specific length with `max_length` (e.g. `max_length=45`) or leave max_length to None to pad to the maximal input size of the model (e.g. 512 for Bert).

FutureWarning,

Original: I was wondering if anyone out there could enlighten me on this car I saw the other day. It was a 2-door sports car, looked to be from the late 60s/early 70s. It was called a Bricklin. The doors were really small. In addition, the front bumper was separate from the rest of the body. This is all I know. If anyone can tell me a model name, engine specs, years of production, where this car is made, history, or whatever info you have on this funky looking car, please e-mail.

Token IDs: torch.Size([512])

```
In [ ]:
```

```
model = BertForSequenceClassification.from_pretrained(
    "bert-base-uncased",
    num_labels = 20,
    output_attentions = False,
    output_hidden_states = False,
)
model.cuda();
```

Some weights of the model checkpoint at bert-base-uncased were not used when initializing BertForSequenceClassification: ['cls.predictions.bias', 'cls.predictions.transform.dense.weight', 'cls.predictions.transform.dense.bias', 'cls.predictions.decoder.weight', 'cls.seq_relationship.weight', 'cls.seq_relationship.bias', 'cls.predictions.transform.LayerNorm.weight', 'cls.predictions.transform.LayerNorm.bias']

- This IS expected if you are initializing BertForSequenceClassification from the checkpoint of a model trained on another task or with another architecture (e.g. initializing a BertForSequenceClassification model from a BertForPreTraining model).

- This IS NOT expected if you are initializing BertForSequenceClassification from the checkpoint of a model that you expect to be exactly identical (initializing a BertForSequenceClassification model from a BertForSequenceClassification model).

Some weights of BertForSequenceClassification were not initialized from the model checkpoint at bert-base-uncased and are newly initialized: ['classifier.weight', 'classifier.bias']

You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.

In []:

```
#Parameters

for param in model.parameters():
    param.requires_grad = False
model.classifier.weight.requires_grad = True
model.classifier.bias.requires_grad = True
```

In []:

```
params['epochs'], params['optimizer'], params['scheduler'] = model_parameters(model)
```

In []:

```
model.to(device);
```

In []:

```
#Train Loss and Validation loss
train_losses, valid_losses = train_loop(model, train_dataloader, validation_dataloader, params['optimizer'], params['scheduler'], params['epochs'], 1234, device)
```

```
Epoch 1/2, Train Loss: 2.8773    Valid Loss: 2.7825, Duration: 0:12:14.064195
Epoch 2/2, Train Loss: 2.7270    Valid Loss: 2.6938, Duration: 0:12:19.591873
```

In []:

```
#Getting accuracies
train_acc = get_accuracy(train_dataloader, model, device)
valid_acc = get_accuracy(validation_dataloader, model, device)
print("Train acc: ", train_acc, "Valid acc: ", valid_acc)
```

```
Train acc:  0.25799893936715573 Valid acc:  0.2457514604354753
```

BERT Model with 128 tokens

In []:

```
params['max_len'] = 128
```

In []:

```
train_dataloader, validation_dataloader = loaders(sentences_train, labels_train, sentences_test, labels_test, params['max_len'], params['batch_size'])
```

```
/usr/local/lib/python3.6/dist-packages/transformers/tokenization_utils_base.py:2022:
FutureWarning: The `pad_to_max_length` argument is deprecated and will be removed in a future version, use `padding=True` or `padding='longest'` to pad to the longest sequence in the batch, or use `padding='max_length'` to pad to a max length. In this case, you can give a specific length with `max_length` (e.g. `max_length=45`) or leave max_length to None to pad to the maximal input size of the model (e.g. 512 for Bert).
  FutureWarning,
```

Original: I was wondering if anyone out there could enlighten me on this car I saw the other day. It was a 2-door sports car, looked to be from the late 60s/early 70s. It was called a Bricklin. The doors were really small. In addition, the front bumper was separate from the rest of the body. This is all I know. If anyone can tell me a model name, engine specs, years of production, where this car is made, history, or whatever info you have on this funky looking car, please e-mail.
Token IDs: torch.Size([128])

In []:

```
model = BertForSequenceClassification.from_pretrained(
```



```

"bert-base-uncased",
num_labels = 20,
output_attentions = False,
output_hidden_states = False,
)
model.cuda();

```

Some weights of the model checkpoint at bert-base-uncased were not used when initializing BertForSequenceClassification: ['cls.predictions.bias', 'cls.predictions.transform.dense.weight', 'cls.predictions.transform.dense.bias', 'cls.predictions.decoder.weight', 'cls.seq_relationship.weight', 'cls.seq_relationship.bias', 'cls.predictions.transform.LayerNorm.weight', 'cls.predictions.transform.LayerNorm.bias']

- This IS expected if you are initializing BertForSequenceClassification from the checkpoint of a model trained on another task or with another architecture (e.g. initializing a BertForSequenceClassification model from a BertForPreTraining model).
- This IS NOT expected if you are initializing BertForSequenceClassification from the checkpoint of a model that you expect to be exactly identical (initializing a BertForSequenceClassification model from a BertForSequenceClassification model).

Some weights of BertForSequenceClassification were not initialized from the model checkpoint at bert-base-uncased and are newly initialized: ['classifier.weight', 'classifier.bias']

You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.

In []:

```

for param in model.parameters():
    param.requires_grad = False
model.classifier.weight.requires_grad = True
model.classifier.bias.requires_grad = True

```

In []:

```

params['epochs'], params['optimizer'], params['scheduler'] = model_parameters(model)

```

In []:

```

#Training
model.to(device);

```

In []:

```

train_losses, valid_losses = train_loop(model, train_dataloader, validation_dataloader, params['optimizer'], params['scheduler'], params['epochs'], 1234, device)

```

```

Epoch 1/2, Train Loss: 2.8839    Valid Loss: 2.7787, Duration: 0:02:43.894216
Epoch 2/2, Train Loss: 2.7245    Valid Loss: 2.6867, Duration: 0:02:43.378686

```

In []:

```

train_acc = get_accuracy(train_dataloader, model, device)
valid_acc = get_accuracy(validation_dataloader, model, device)
print(train_acc, valid_acc)

```

```

0.2647162807141594 0.2624800849707913

```

BERT Model with 140 tokens

In []:

```

params['max_len'] = 140

```

In []:

```

train_dataloader, validation_dataloader = loaders(sentences_train, labels_train, sentences_test, labels_test, params['max_len'], params['batch_size'])

```

/usr/local/lib/python3.6/dist-packages/transformers/tokenization_utils_base.py:2022:

```
../usr/local/lib/python3.9/dist-packages/transformers/tokenization_utils_base.py:2022:
FutureWarning: The `pad_to_max_length` argument is deprecated and will be removed in a future vers
ion, use `padding=True` or `padding='longest'` to pad to the longest sequence in the batch, or use
`padding='max_length'` to pad to a max length. In this case, you can give a specific length with `
max_length` (e.g. `max_length=45`) or leave max_length to None to pad to the maximal input size of
the model (e.g. 512 for Bert).
FutureWarning,
```

Original: I was wondering if anyone out there could enlighten me on this car I saw the other day. It was a 2-door sports car, looked to be from the late 60s/early 70s. It was called a Bricklin. The doors were really small. In addition, the front bumper was separate from the rest of the body. This is all I know. If anyone can tell me a model name, engine specs, years of production, where this car is made, history, or whatever info you have on this funky looking car, please e-mail.
Token IDs: torch.Size([140])

In []:

```
model = BertForSequenceClassification.from_pretrained(
    "bert-base-uncased",
    num_labels = 20,
    output_attentions = False,
    output_hidden_states = False,
)
model.cuda();
```

Some weights of the model checkpoint at bert-base-uncased were not used when initializing BertForSequenceClassification: ['cls.predictions.bias', 'cls.predictions.transform.dense.weight', 'cls.predictions.transform.dense.bias', 'cls.predictions.decoder.weight', 'cls.seq_relationship.weight', 'cls.seq_relationship.bias', 'cls.predictions.transform.LayerNorm.weight', 'cls.predictions.transform.LayerNorm.bias'] - This IS expected if you are initializing BertForSequenceClassification from the checkpoint of a model trained on another task or with another architecture (e.g. initializing a BertForSequenceClassification model from a BertForPreTraining model). - This IS NOT expected if you are initializing BertForSequenceClassification from the checkpoint of a model that you expect to be exactly identical (initializing a BertForSequenceClassification model from a BertForSequenceClassification model). Some weights of BertForSequenceClassification were not initialized from the model checkpoint at bert-base-uncased and are newly initialized: ['classifier.weight', 'classifier.bias'] You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.

In []:

```
for param in model.parameters():
    param.requires_grad = False
model.classifier.weight.requires_grad = True
model.classifier.bias.requires_grad = True
```

In []:

```
params['epochs'], params['optimizer'], params['scheduler'] = model_parameters(model)
```

In []:

```
train_losses, valid_losses = train_loop(model, train_dataloader, validation_dataloader, params['optimizer'], params['scheduler'], params['epochs'], 1234, device)
```

```
Epoch 1/2, Train Loss: 2.8841    Valid Loss: 2.7832, Duration: 0:03:09.593033
Epoch 2/2, Train Loss: 2.7236    Valid Loss: 2.6880, Duration: 0:03:08.959613
```

In []:

```
train_acc = get_accuracy(train_dataloader, model, device)
valid_acc = get_accuracy(validation_dataloader, model, device)
print(train_acc, valid_acc)
```

```
0.2672794767544635 0.2627456186935741
```

BERT Model 4 by truncating from the end

In []:

```
def tokenizer_s(sentences, max_len):
    input_ids = []
    attention_masks = []
    for sent in sentences:
        sent = sent.split(" ")
        sent.reverse()
        sent = " ".join(sent)
        encoded_dict = tokenizer.encode_plus(
            sent,
            add_special_tokens = True,
            max_length = max_len,
            pad_to_max_length = True,
            return_attention_mask = True,
            return_tensors = 'pt',
        )

        # Add the encoded sentence to the list.
        input_ids.append(encoded_dict['input_ids'])

        # And its attention mask (simply differentiates padding from non-padding).
        attention_masks.append(encoded_dict['attention_mask'])

    # Convert the lists into tensors.
    input_ids = torch.cat(input_ids, dim=0)
    attention_masks = torch.cat(attention_masks, dim=0)

    return input_ids, attention_masks
```

In []:

```
params['max_len'] = 140
```

In []:

```
train_dataloader, validation_dataloader = loaders(sentences_train, labels_train, sentences_test, labels_test, params['max_len'], params['batch_size'])
```

```
/usr/local/lib/python3.6/dist-packages/transformers/tokenization_utils_base.py:2022:
FutureWarning: The `pad_to_max_length` argument is deprecated and will be removed in a future version, use `padding=True` or `padding='longest'` to pad to the longest sequence in the batch, or use `padding='max_length'` to pad to a max length. In this case, you can give a specific length with `max_length` (e.g. `max_length=45`) or leave max_length to None to pad to the maximal input size of the model (e.g. 512 for Bert).
FutureWarning,
```

Original: I was wondering if anyone out there could enlighten me on this car I saw the other day. It was a 2-door sports car, looked to be from the late 60s/early 70s. It was called a Bricklin. The doors were really small. In addition, the front bumper was separate from the rest of the body. This is all I know. If anyone can tell me a model name, engine specs, years of production, where this car is made, history, or whatever info you have on this funky looking car, please e-mail.
Token IDs: torch.Size([140])

In []:

```
model = BertForSequenceClassification.from_pretrained(
    "bert-base-uncased",
    num_labels = 20,
    output_attentions = False,
    output_hidden_states = False,
)
model.cuda();
```

```
Some weights of the model checkpoint at bert-base-uncased were not used when initializing BertForSequenceClassification: ['cls.predictions.bias', 'cls.predictions.transform.dense.weight',
```

```

cls.predictions.transform.dense.bias', 'cls.predictions.decoder.weight',
'cls.seq_relationship.weight', 'cls.seq_relationship.bias',
'cls.predictions.transform.LayerNorm.weight', 'cls.predictions.transform.LayerNorm.bias']
- This IS expected if you are initializing BertForSequenceClassification from the checkpoint of a
model trained on another task or with another architecture (e.g. initializing a
BertForSequenceClassification model from a BertForPreTraining model).
- This IS NOT expected if you are initializing BertForSequenceClassification from the checkpoint o
f a model that you expect to be exactly identical (initializing a BertForSequenceClassification mo
del from a BertForSequenceClassification model).
Some weights of BertForSequenceClassification were not initialized from the model checkpoint at be
rt-base-uncased and are newly initialized: ['classifier.weight', 'classifier.bias']
You should probably TRAIN this model on a down-stream task to be able to use it for predictions an
d inference.

```

In []:

```

for param in model.parameters():
    param.requires_grad = False
model.classifier.weight.requires_grad = True
model.classifier.bias.requires_grad = True

```

In []:

```

params['epochs'], params['optimizer'], params['scheduler'] = model_parameters(model)

```

In []:

```

train_losses, valid_losses = train_loop(model, train_dataloader, validation_dataloader, params['opt
imizer'], params['scheduler'], params['epochs'], 1234, device)

```

```

Epoch 1/2, Train Loss: 2.8720    Valid Loss: 2.7455, Duration: 0:03:09.559264
Epoch 2/2, Train Loss: 2.7120    Valid Loss: 2.6613, Duration: 0:03:08.964845

```

In []:

```

train_acc = get_accuracy(train_dataloader, model, device)
valid_acc = get_accuracy(validation_dataloader, model, device)
print("Train acc:", train_acc, "Valid acc:", valid_acc)

```

```

Train acc: 0.27346650167933534 Valid acc: 0.2705788635156665

```

For Sequence length, 128

In []:

```

params['max_len'] = 128

```

In []:

```

train_dataloader, validation_dataloader = loaders(sentences_train, labels_train, sentences_test, la
bels_test, params['max_len'], params['batch_size'])

```

```

/usr/local/lib/python3.6/dist-packages/transformers/tokenization_utils_base.py:2022:
FutureWarning: The `pad_to_max_length` argument is deprecated and will be removed in a future vers
ion, use `padding=True` or `padding='longest'` to pad to the longest sequence in the batch, or use
`padding='max_length'` to pad to a max length. In this case, you can give a specific length with `
max_length` (e.g. `max_length=45`) or leave max_length to None to pad to the maximal input size of
the model (e.g. 512 for Bert).
FutureWarning,

```

Original: I was wondering if anyone out there could enlighten me on this car I saw the other day. It was a 2-door sports car, looked to be from the late 60s/early 70s. It was called a Bricklin. The doors were really small. In addition, the front bumper was separate from the rest of the body. This is all I know. If anyone can tell me a model name, engine specs, years of production, where this car is made, history, or whatever info you have on this funky looking car, please e-mail.

Token IDs: torch.Size([128])

In []:

```
model = BertForSequenceClassification.from_pretrained(
    "bert-base-uncased",
    num_labels = 20,
    output_attentions = False,
    output_hidden_states = False,
)
model.cuda();
```

Some weights of the model checkpoint at bert-base-uncased were not used when initializing BertForSequenceClassification: ['cls.predictions.bias', 'cls.predictions.transform.dense.weight', 'cls.predictions.transform.dense.bias', 'cls.predictions.decoder.weight', 'cls.seq_relationship.weight', 'cls.seq_relationship.bias', 'cls.predictions.transform.LayerNorm.weight', 'cls.predictions.transform.LayerNorm.bias']

- This IS expected if you are initializing BertForSequenceClassification from the checkpoint of a model trained on another task or with another architecture (e.g. initializing a BertForSequenceClassification model from a BertForPreTraining model).
- This IS NOT expected if you are initializing BertForSequenceClassification from the checkpoint of a model that you expect to be exactly identical (initializing a BertForSequenceClassification model from a BertForSequenceClassification model).

Some weights of BertForSequenceClassification were not initialized from the model checkpoint at bert-base-uncased and are newly initialized: ['classifier.weight', 'classifier.bias']

You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.

In []:

```
for param in model.parameters():
    param.requires_grad = False
model.classifier.weight.requires_grad = True
model.classifier.bias.requires_grad = True
```

In []:

```
params['epochs'], params['optimizer'], params['scheduler'] = model_parameters(model)
```

In []:

```
#Train valid loss
train_losses, valid_losses = train_loop(model, train_dataloader, validation_dataloader, params['optimizer'], params['scheduler'], params['epochs'], 1234, device)
```

Epoch 1/2, Train Loss: 2.8730 Valid Loss: 2.7592, Duration: 0:02:41.486860
Epoch 2/2, Train Loss: 2.7193 Valid Loss: 2.6702, Duration: 0:02:42.259626

In []:

```
#Train valid accuracy
train_acc = get_accuracy(train_dataloader, model, device)
valid_acc = get_accuracy(validation_dataloader, model, device)
print("Train acc:", train_acc, "Valid acc:", valid_acc)
```

Train acc: 0.27196393848329503 Valid acc: 0.2586298459904408

For Sequence length, 512

In []:

```
params['max_len'] = 512
```

In []:

```
train_dataloader, validation_dataloader = loaders(sentences_train, labels_train, sentences_test, labels_test, params['max_len'], params['batch_size'])
```

```
2020-0000, paramet_max_len=1, paramet_max_len=1,
```

```
/usr/local/lib/python3.6/dist-packages/transformers/tokenization_utils_base.py:2022:
FutureWarning: The `pad_to_max_length` argument is deprecated and will be removed in a future vers
ion, use `padding=True` or `padding='longest'` to pad to the longest sequence in the batch, or use
`padding='max_length'` to pad to a max length. In this case, you can give a specific length with `
max_length` (e.g. `max_length=45`) or leave max_length to None to pad to the maximal input size of
the model (e.g. 512 for Bert).
FutureWarning,
```

Original: I was wondering if anyone out there could enlighten me on this car I saw the other day. It was a 2-door sports car, looked to be from the late 60s/early 70s. It was called a Bricklin. The doors were really small. In addition, the front bumper was separate from the rest of the body. This is all I know. If anyone can tell me a model name, engine specs, years of production, where this car is made, history, or whatever info you have on this funky looking car, please e-mail.
Token IDs: torch.Size([512])

In []:

```
#Model
model = BertForSequenceClassification.from_pretrained(
    "bert-base-uncased",
    num_labels = 20,
    output_attentions = False,
    output_hidden_states = False,
)
model.cuda();
```

Some weights of the model checkpoint at bert-base-uncased were not used when initializing BertForSequenceClassification: ['cls.predictions.bias', 'cls.predictions.transform.dense.weight', 'cls.predictions.transform.dense.bias', 'cls.predictions.decoder.weight', 'cls.seq_relationship.weight', 'cls.seq_relationship.bias', 'cls.predictions.transform.LayerNorm.weight', 'cls.predictions.transform.LayerNorm.bias']
- This IS expected if you are initializing BertForSequenceClassification from the checkpoint of a model trained on another task or with another architecture (e.g. initializing a BertForSequenceClassification model from a BertForPreTraining model).
- This IS NOT expected if you are initializing BertForSequenceClassification from the checkpoint of a model that you expect to be exactly identical (initializing a BertForSequenceClassification model from a BertForSequenceClassification model).
Some weights of BertForSequenceClassification were not initialized from the model checkpoint at bert-base-uncased and are newly initialized: ['classifier.weight', 'classifier.bias']
You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.

In []:

```
for param in model.parameters():
    param.requires_grad = False
model.classifier.weight.requires_grad = True
model.classifier.bias.requires_grad = True
```

In []:

```
params['epochs'], params['optimizer'], params['scheduler'] = model_parameters(model)
```

In []:

```
#Train valid loss
train_losses, valid_losses = train_loop(model, train_dataloader, validation_dataloader, params['optimizer'], params['scheduler'], params['epochs'], 1234, device)
```

```
Epoch 1/2, Train Loss: 2.8569    Valid Loss: 2.7333, Duration: 0:12:16.321910
Epoch 2/2, Train Loss: 2.6893    Valid Loss: 2.6386, Duration: 0:12:17.209193
```

In []:

```
#train valid accuracy
train_acc = get_accuracy(train_dataloader, model, device)
```

```
valid_acc = get_accuracy(validation_dataloader, model, device)
print("Train acc:", train_acc, "Valid acc:", valid_acc)
```

Train acc: 0.2864592540215662 Valid acc: 0.2831917153478492

BERT Model - Truncating from the mid

In []:

```
def tokenizer_s(sentences, max_len):
    input_ids = []
    attention_masks = []
    for sent in sentences:
        sent = sent.split(" ")
        sent.reverse()
        sent = " ".join(sent)
        encoded_dict = tokenizer.encode_plus(
            sent,
            add_special_tokens = True,
            max_length = max_len,
            pad_to_max_length = True,
            return_attention_mask = True,
            return_tensors = 'pt',
        )

        # Add the encoded sentence to the list.
        input_ids.append(encoded_dict['input_ids'])

        # And its attention mask (simply differentiates padding from non-padding).
        attention_masks.append(encoded_dict['attention_mask'])

    # Convert the lists into tensors.
    input_ids = torch.cat(input_ids, dim=0)
    attention_masks = torch.cat(attention_masks, dim=0)

    return input_ids, attention_masks
```

For Sequence length, 140

In []:

```
params['max_len'] = 140
```

In []:

```
train_dataloader, validation_dataloader = loaders(sentences_train, labels_train, sentences_test, labels_test, params['max_len'], params['batch_size'])
```

```
/usr/local/lib/python3.6/dist-packages/transformers/tokenization_utils_base.py:2022:
FutureWarning: The `pad_to_max_length` argument is deprecated and will be removed in a future version, use `padding=True` or `padding='longest'` to pad to the longest sequence in the batch, or use `padding='max_length'` to pad to a max length. In this case, you can give a specific length with `max_length` (e.g. `max_length=45`) or leave max_length to None to pad to the maximal input size of the model (e.g. 512 for Bert).
FutureWarning,
```

Original: I was wondering if anyone out there could enlighten me on this car I saw the other day. It was a 2-door sports car, looked to be from the late 60s/early 70s. It was called a Bricklin. The doors were really small. In addition, the front bumper was separate from the rest of the body. This is all I know. If anyone can tell me a model name, engine specs, years of production, where this car is made, history, or whatever info you have on this funky looking car, please e-mail.
Token IDs: torch.Size([140])

In []:

```
model = BertForSequenceClassification.from_pretrained(
    "bert-base-uncased",
```

```

num_labels = 20,
output_attentions = False,
output_hidden_states = False,
)
model.cuda();

```

Some weights of the model checkpoint at bert-base-uncased were not used when initializing BertForSequenceClassification: ['cls.predictions.bias', 'cls.predictions.transform.dense.weight', 'cls.predictions.transform.dense.bias', 'cls.predictions.decoder.weight', 'cls.seq_relationship.weight', 'cls.seq_relationship.bias', 'cls.predictions.transform.LayerNorm.weight', 'cls.predictions.transform.LayerNorm.bias']

- This IS expected if you are initializing BertForSequenceClassification from the checkpoint of a model trained on another task or with another architecture (e.g. initializing a BertForSequenceClassification model from a BertForPreTraining model).
- This IS NOT expected if you are initializing BertForSequenceClassification from the checkpoint of a model that you expect to be exactly identical (initializing a BertForSequenceClassification model from a BertForSequenceClassification model).

Some weights of BertForSequenceClassification were not initialized from the model checkpoint at bert-base-uncased and are newly initialized: ['classifier.weight', 'classifier.bias']

You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.

In []:

```

for param in model.parameters():
    param.requires_grad = False
model.classifier.weight.requires_grad = True
model.classifier.bias.requires_grad = True

```

In []:

```

params['epochs'], params['optimizer'], params['scheduler'] = model_parameters(model)

```

In []:

```

#Train valid loss
train_losses, valid_losses = train_loop(model, train_dataloader, validation_dataloader, params['optimizer'], params['scheduler'], params['epochs'], 1234, device)

```

```

Epoch 1/2, Train Loss: 2.8720    Valid Loss: 2.7455, Duration: 0:03:02.680551
Epoch 2/2, Train Loss: 2.7120    Valid Loss: 2.6613, Duration: 0:03:09.193130

```

In []:

```

#Train valid accuracy
train_acc = get_accuracy(train_dataloader, model, device)
valid_acc = get_accuracy(validation_dataloader, model, device)
print("Train accuracy:", train_acc, "Valid accuracy:", valid_acc)

```

```

Train accuracy: 0.27346650167933534 Valid accuracy: 0.2705788635156665

```

Chunking

In []:

```

def chunking(sentences, labels, max_length):
    final_sentences = []
    final_labels = []
    order = []
    for target_order, (sent, label) in enumerate(zip(sentences, labels)):
        sent = sent.split(" ")
        length = len(sent)
        parts = math.ceil(length/max_length)
        for i in range(1, parts + 1):
            chunk = " ".join(sent[max_length*(i-1) : max_length*i])
            final_sentences.append(chunk)
            final_labels.append(label)
            order.append(target_order)
    # return final_sentences, final_labels, order

```



```
# return final_sentences, final_labels, order
return np.array(final_sentences), np.array(final_labels), np.array(order)
```

In []:

```
params['max_length'] = 140
sentences_train_chunked, labels_train_chunked, order_train = chunking(sentences_train,
labels_train, params['max_length'])
sentences_test_chunked, labels_test_chunked, order_test = chunking(sentences_test, labels_test, par
ams['max_length'])
```

In []:

```
#Creating the model
def tokenize_s(sentences, max_len):
    input_ids = []
    attention_masks = []
    for sent in sentences:
        sent = sent.split(" ")
        sent = preprocessing(sent, max_len)
        encoded_dict = tokenizer.encode_plus(
            sent,                                # Sentence to encode.
            add_special_tokens = True, # Add '[CLS]' and '[SEP]'
            max_length = max_len,           # Pad & truncate all sentences.
            pad_to_max_length = True,
            return_attention_mask = True,   # Construct attn. masks.
            return_tensors = 'pt',         # Return pytorch tensors.
        )

        # Add the encoded sentence to the list.
        input_ids.append(encoded_dict['input_ids'])

        # And its attention mask (simply differentiates padding from non-padding).
        attention_masks.append(encoded_dict['attention_mask'])

    # Convert the lists into tensors.
    input_ids = torch.cat(input_ids, dim=0)
    attention_masks = torch.cat(attention_masks, dim=0)

    return input_ids, attention_masks
```

In []:

```
def processing(sent, max_len):
    return sent
```

For sequence length, 140

In []:

```
params['max_len'] = 140
train_dataloader, validation_dataloader = loaders(sentences_train_chunked, labels_train_chunked,
sentences_test_chunked,                                labels_test_chunked, params['max_len'], para
['batch_size'])
```

Truncation was not explicitly activated but `max_length` is provided a specific value, please use `truncation=True` to explicitly truncate examples to max length. Defaulting to 'longest_first' truncation strategy. If you encode pairs of sequences (GLUE-style) with the tokenizer you can select this strategy more precisely by providing a specific strategy to `truncation`.

/usr/local/lib/python3.6/dist-packages/transformers/tokenization_utils_base.py:2022: FutureWarning: The `pad_to_max_length` argument is deprecated and will be removed in a future version, use `padding=True` or `padding='longest'` to pad to the longest sequence in the batch, or use `padding='max_length'` to pad to a max length. In this case, you can give a specific length with `max_length` (e.g. `max_length=45`) or leave max_length to None to pad to the maximal input size of the model (e.g. 512 for Bert).

FutureWarning,

Original: I was wondering if anyone out there could enlighten me on this car I saw the other day. It was a 2-door sports car, looked to be from the late 60s/early 70s. It was called a Bricklin. The doors were really small. In addition,

the front bumper was separate from the rest of the body. This is all I know. If anyone can tell me a model name, engine specs, years of production, where this car is made, history, or whatever info you have on this funky looking car, please e-mail.

Token IDs: torch.Size([140])

In []:

```
model = BertForSequenceClassification.from_pretrained(
    "bert-base-uncased",
    num_labels = 20,
    output_attentions = False,
    output_hidden_states = False,
)
model.cuda();
```

Some weights of the model checkpoint at bert-base-uncased were not used when initializing BertForSequenceClassification: ['cls.predictions.bias', 'cls.predictions.transform.dense.weight', 'cls.predictions.transform.dense.bias', 'cls.predictions.decoder.weight', 'cls.seq_relationship.weight', 'cls.seq_relationship.bias', 'cls.predictions.transform.LayerNorm.weight', 'cls.predictions.transform.LayerNorm.bias']

- This IS expected if you are initializing BertForSequenceClassification from the checkpoint of a model trained on another task or with another architecture (e.g. initializing a BertForSequenceClassification model from a BertForPreTraining model).
- This IS NOT expected if you are initializing BertForSequenceClassification from the checkpoint of a model that you expect to be exactly identical (initializing a BertForSequenceClassification model from a BertForSequenceClassification model).

Some weights of BertForSequenceClassification were not initialized from the model checkpoint at bert-base-uncased and are newly initialized: ['classifier.weight', 'classifier.bias']

You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.

In []:

```
for param in model.parameters():
    param.requires_grad = False
model.classifier.weight.requires_grad = True
model.classifier.bias.requires_grad = True
```

In []:

```
params['epochs'], params['optimizer'], params['scheduler'] = model_parameters(model)
```

In []:

```
train_losses, valid_losses = train_loop(model, train_dataloader, validation_dataloader, params['optimizer'], params['scheduler'], params['epochs'], 1234, device)
```

```
Epoch 1/2, Train Loss: 2.9801    Valid Loss: 2.9615, Duration: 0:05:35.546815
Epoch 2/2, Train Loss: 2.9582    Valid Loss: 2.9563, Duration: 0:05:46.537441
```

In []:

```
def get_chunk_accuracy(model, data_iter, order, device):
    predictions = []
    correct_labels = []

    model.eval()
    with torch.no_grad():

        for batch in data_iter:

            b_input_ids = batch[0].to(device)
            b_input_mask = batch[1].to(device)
            b_labels = batch[2].to(device)
            # forward pass
```

```

    loss, logits = model(b_input_ids,
                        token_type_ids=None,
                        attention_mask=b_input_mask,
                        labels=b_labels)

    _, indices = torch.max(logits, dim=1)

    indices = indices.cpu().numpy()
    predictions += list(indices)
    b_labels = b_labels.cpu().numpy()
    correct_labels += list(b_labels)

    return predictions, correct_labels, order

```

In []:

```

predictions, correct_labels, order = get_chunk_accuracy(model, validation_dataloader, order_test, device)

```

In []:

```

state = pd.DataFrame({'order':order, 'preds':predictions, 'correct':correct_labels})
state_2 = state.groupby('order').agg(lambda x:x.value_counts().index[0])
state_2[state_2['preds']==state_2['correct']].shape[0]/state['order'].nunique()

```

Out[]:

```

0.06120552310143388

```

In []:

```

del sentences_train_chunked, labels_train_chunked, order_train, sentences_test_chunked,
labels_test_chunked, order_test

```

In []:

```

import gc
gc.collect()

```

Out[]:

```

222

```

In experiment 1, after reducing the sequence length from 512 we observe 1% improvement in the accuracy of the model whereas 'truncating the model from end' & 'truncating the model from mid' gives us the same accuracy as the one observed after reducing the sequence length. In case of chunking we observe significant reduction of 20% in comparison to other model.

TASK 2 - Fine tune Bert

In []:

```

def tokenize_s(sentences, max_len):
    input_ids = []
    attention_masks = []
    for sent in sentences:
        sent = sent.split(" ")
        sent.reverse()
        sent = " ".join(sent)
        encoded_dict = tokenizer.encode_plus(
            sent,
            add_special_tokens = True,
            max_length = max_len,
            pad_to_max_length = True,
            return_attention_mask = True,
            return_tensors = 'pt',
        )

    # Add the encoded sentence to the list.

```

```

input_ids.append(encoded_dict['input_ids'])

# And its attention mask (simply differentiates padding from non-padding).
attention_masks.append(encoded_dict['attention_mask'])

# Convert the lists into tensors.
input_ids = torch.cat(input_ids, dim=0)
attention_masks = torch.cat(attention_masks, dim=0)

return input_ids, attention_masks

```

Sequence Length, 512

In []:

```

params['max_len'] = 512
params['batch_size'] = 8

```

In []:

```

train_dataloader, validation_dataloader = loaders(sentences_train, labels_train, sentences_test, labels_test, params['max_len'], params['batch_size'])

```

/usr/local/lib/python3.6/dist-packages/transformers/tokenization_utils_base.py:2022: FutureWarning: The `pad_to_max_length` argument is deprecated and will be removed in a future version, use `padding=True` or `padding='longest'` to pad to the longest sequence in the batch, or use `padding='max_length'` to pad to a max length. In this case, you can give a specific length with `max_length` (e.g. `max_length=45`) or leave max_length to None to pad to the maximal input size of the model (e.g. 512 for Bert).

FutureWarning,

Original: I was wondering if anyone out there could enlighten me on this car I saw the other day. It was a 2-door sports car, looked to be from the late 60s/early 70s. It was called a Bricklin. The doors were really small. In addition, the front bumper was separate from the rest of the body. This is all I know. If anyone can tell me a model name, engine specs, years of production, where this car is made, history, or whatever info you have on this funky looking car, please e-mail.

Token IDs: torch.Size([512])

In []:

```

model = BertForSequenceClassification.from_pretrained(
    "bert-base-uncased",
    num_labels = 20,
    output_attentions = False,
    output_hidden_states = False,
)
model.cuda();

```

Some weights of the model checkpoint at bert-base-uncased were not used when initializing BertForSequenceClassification: ['cls.predictions.bias', 'cls.predictions.transform.dense.weight', 'cls.predictions.transform.dense.bias', 'cls.predictions.decoder.weight', 'cls.seq_relationship.weight', 'cls.seq_relationship.bias', 'cls.predictions.transform.LayerNorm.weight', 'cls.predictions.transform.LayerNorm.bias']

- This IS expected if you are initializing BertForSequenceClassification from the checkpoint of a model trained on another task or with another architecture (e.g. initializing a BertForSequenceClassification model from a BertForPreTraining model).
- This IS NOT expected if you are initializing BertForSequenceClassification from the checkpoint of a model that you expect to be exactly identical (initializing a BertForSequenceClassification model from a BertForSequenceClassification model).

Some weights of BertForSequenceClassification were not initialized from the model checkpoint at bert-base-uncased and are newly initialized: ['classifier.weight', 'classifier.bias']

You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.

In []:

```

params['epochs'], params['optimizer'], params['scheduler'] = model_parameters(model)
torch.cuda.empty_cache()

```

```
print(torch.cuda.memory_summary(device=None, abbreviated=False))
```

```
=====
|
|                PyTorch CUDA memory summary, device ID 0
|
|-----|
|          CUDA OOMs: 0          |          cudaMalloc retries: 0          | | | |
|---|---|---|---|---|
|          Metric          | Cur Usage | Peak Usage | Tot Alloc | Tot Freed |
|-----|
| Allocated memory        |      858 MB |      7938 MB | 107110 GB | 107109 GB |
|   from large pool       |      857 MB |      7936 MB | 107088 GB | 107087 GB |
|   from small pool       |       1 MB |       21 MB |    22 GB |    22 GB |
|-----|
| Active memory           |      858 MB |      7938 MB | 107110 GB | 107109 GB |
|   from large pool       |      857 MB |      7936 MB | 107088 GB | 107087 GB |
|   from small pool       |       1 MB |       21 MB |    22 GB |    22 GB |
|-----|
| GPU reserved memory     |      924 MB |      8234 MB |   8664 MB |   7740 MB |
|   from large pool       |      922 MB |      8206 MB |   8636 MB |   7714 MB |
|   from small pool       |       2 MB |       28 MB |    28 MB |    26 MB |
|-----|
| Non-releasable memory   |    67086 KB |   261595 KB |   27587 GB |   27587 GB |
|   from large pool       |    66130 KB |   259666 KB |   27564 GB |   27564 GB |
|   from small pool       |     956 KB |     7075 KB |    22 GB |    22 GB |
|-----|
| Allocations             |      405   |      1066   |    8230 K |    8229 K |
|   from large pool       |      150   |       496   |    5096 K |    5096 K |
|   from small pool       |      255   |       714   |    3133 K |    3133 K |
|-----|
| Active allocs           |      405   |      1066   |    8230 K |    8229 K |
|   from large pool       |      150   |       496   |    5096 K |    5096 K |
|   from small pool       |      255   |       714   |    3133 K |    3133 K |
|-----|
| GPU reserved segments   |       51   |       246   |     264   |     213   |
|   from large pool       |       50   |       232   |     250   |     200   |
|   from small pool       |        1   |        14   |      14   |      13   |
|-----|
| Non-releasable allocs   |       44   |        59   |    3482 K |    3482 K |
|   from large pool       |       25   |        33   |    2465 K |    2465 K |
|   from small pool       |       19   |        31   |    1016 K |    1016 K |
|-----|
|=====
```

```
In [ ]:
```

```
train_losses, valid_losses = train_loop(model, train_dataloader, validation_dataloader, params['optimizer'], params['scheduler'], params['epochs'],1234, device)
```

```
Epoch 1/2, Train Loss: 1.0224      Valid Loss: 1.1226, Duration: 0:23:04.048650
Epoch 2/2, Train Loss: 1.0085      Valid Loss: 1.1113, Duration: 0:23:30.232013
```

```
In [ ]:
```

```
train_acc = get_accuracy(train_dataloader, model, device)
valid_acc = get_accuracy(validation_dataloader, model, device)
print("Train accuracy:", train_acc, "Valid Accuracy:" ,valid_acc)
```

```
Train accuracy: 0.782440232714563428 Valid Accuracy: 0.6912464132147456348
```

After fine-tuning the model we observe significant improvement in both training and validation accuracy.

TASK 3 - Different BERT variations from Huggingface library

ALBERT

```
In [ ]:
```

```
from transformers import AlbertTokenizer, AlbertForSequenceClassification
```

```
In [ ]:
```

```
tokenizer = AlbertTokenizer.from_pretrained('albert-base-v2')
```

```
In [ ]:
```

```
def tokenize_s(sentences, max_len):
    input_ids = []
    attention_masks = []
    for sent in sentences:
        sent = sent.split(" ")
        sent.reverse()
        sent = " ".join(sent)
        encoded_dict = tokenizer.encode_plus(
            sent,                                # Sentence to encode.
            add_special_tokens = True,           # Add '[CLS]' and '[SEP]'
            max_length = max_len,                # Pad & truncate all sentences.
            pad_to_max_length = True,
            return_attention_mask = True,        # Construct attn. masks.
            return_tensors = 'pt',              # Return pytorch tensors.
        )

        # Add the encoded sentence to the list.
        input_ids.append(encoded_dict['input_ids'])

        # And its attention mask (simply differentiates padding from non-padding).
        attention_masks.append(encoded_dict['attention_mask'])

    # Convert the lists into tensors.
    input_ids = torch.cat(input_ids, dim=0)
    attention_masks = torch.cat(attention_masks, dim=0)

    return input_ids, attention_masks
```

Sequence Length, 512

```
In [ ]:
```

```
params['max_len'] = 512
params['batch_size'] = 8
```

```
In [ ]:
```

```
train_dataloader, validation_dataloader = loaders(sentences_train, labels_train, sentences_test, labels_test, params['max_len'], params['batch_size'])
```

Truncation was not explicitly activated but `max_length` is provided a specific value, please use `truncation=True` to explicitly truncate examples to max length. Defaulting to 'longest_first' truncation strategy. If you encode pairs of sequences (GLUE-style) with the tokenizer you can select this strategy more precisely by providing a specific strategy to `truncation`.

/usr/local/lib/python3.6/dist-packages/transformers/tokenization_utils_base.py:2022: FutureWarning: The `pad_to_max_length` argument is deprecated and will be removed in a future version, use `padding=True` or `padding='longest'` to pad to the longest sequence in the batch, or use `padding='max_length'` to pad to a max length. In this case, you can give a specific length with `max_length` (e.g. `max_length=45`) or leave max_length to None to pad to the maximal input size of the model (e.g. 512 for Bert).

FutureWarning,

Original: I was wondering if anyone out there could enlighten me on this car I saw the other day. It was a 2-door sports car, looked to be from the late 60s/early 70s. It was called a Bricklin. The doors were really small. In addition, the front bumper was separate from the rest of the body. This is all I know. If anyone can tell me a model name, engine specs, years of production, where this car is made, history, or whatever info you have on this funky looking car, please e-mail.

Token IDs: torch.Size([512])

```
In [ ]:
```

```

model = AlbertForSequenceClassification.from_pretrained(
    "albert-base-v2",
    num_labels = 20,
    output_attentions = False,
    output_hidden_states = False,
)
model.cuda();

```

Some weights of the model checkpoint at albert-base-v2 were not used when initializing AlbertForSequenceClassification: ['predictions.bias', 'predictions.LayerNorm.weight', 'predictions.LayerNorm.bias', 'predictions.dense.weight', 'predictions.dense.bias', 'predictions.decoder.weight', 'predictions.decoder.bias']

- This IS expected if you are initializing AlbertForSequenceClassification from the checkpoint of a model trained on another task or with another architecture (e.g. initializing a BertForSequenceClassification model from a BertForPreTraining model).
- This IS NOT expected if you are initializing AlbertForSequenceClassification from the checkpoint of a model that you expect to be exactly identical (initializing a BertForSequenceClassification model from a BertForSequenceClassification model).

Some weights of AlbertForSequenceClassification were not initialized from the model checkpoint at albert-base-v2 and are newly initialized: ['classifier.weight', 'classifier.bias']

You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.

Without Fine-tuning

In []:

```

model = AlbertForSequenceClassification.from_pretrained(
    "albert-base-v2",
    num_labels = 20,
    output_attentions = False,
    output_hidden_states = False,
)
model.cuda();

```

Some weights of the model checkpoint at albert-base-v2 were not used when initializing AlbertForSequenceClassification: ['predictions.bias', 'predictions.LayerNorm.weight', 'predictions.LayerNorm.bias', 'predictions.dense.weight', 'predictions.dense.bias', 'predictions.decoder.weight', 'predictions.decoder.bias']

- This IS expected if you are initializing AlbertForSequenceClassification from the checkpoint of a model trained on another task or with another architecture (e.g. initializing a BertForSequenceClassification model from a BertForPreTraining model).
- This IS NOT expected if you are initializing AlbertForSequenceClassification from the checkpoint of a model that you expect to be exactly identical (initializing a BertForSequenceClassification model from a BertForSequenceClassification model).

Some weights of AlbertForSequenceClassification were not initialized from the model checkpoint at albert-base-v2 and are newly initialized: ['classifier.weight', 'classifier.bias']

You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.

In []:

```

for param in model.parameters():
    param.requires_grad = False
for param in model.classifier.parameters():
    param.requires_grad = True

```

In []:

```

params['epochs'], params['optimizer'], params['scheduler'] = model_parameters(model)

```

In []:

```

train_losses, valid_losses = train_loop(model, train_dataloader, validation_dataloader, params['optimizer'], params['scheduler'], params['epochs'], 1234, device)

```

```

Epoch 1/2, Train Loss: 2.9722    Valid Loss: 2.9379, Duration: 0:11:57.462271
Epoch 2/2, Train Loss: 2.9219    Valid Loss: 2.9135, Duration: 0:11:58.317175

```

In []:

```
train_acc = get_accuracy(train_dataloader, model, device)
valid_acc = get_accuracy(validation_dataloader, model, device)
print("Train accuracy:", train_acc, "Valid Accuracy:" ,valid_acc)
```

Train accuracy: 0.1251546756231218 Valid Accuracy: 0.12705788635156665

With Fine-Tuning

In []:

```
params['epochs'], params['optimizer'], params['scheduler'] = model_parameters(model)
```

In []:

```
train_losses, valid_losses = train_loop(model, train_dataloader, validation_dataloader, params['optimizer'], params['scheduler'], params['epochs'],1234, device)
```

Epoch 1/2, Train Loss: 1.6314 Valid Loss: 1.2753, Duration: 0:23:53.068745
Epoch 2/2, Train Loss: 0.9262 Valid Loss: 1.1068, Duration: 0:23:54.198061

In []:

```
train_acc = get_accuracy(train_dataloader, model, device)
valid_acc = get_accuracy(validation_dataloader, model, device)
print("Train accuracy:", train_acc, "Valid Accuracy:" ,valid_acc)
```

Train accuracy: 0.7959165635495846 Valid Accuracy: 0.6733935209771641

In case of "ALBERTA" we observe similar results as in case of "bert-base-uncased" for fine-tuned model but in case of without fine-tuned model we observe 12% higher accuracy for "bert-base-uncased" model.

ROBERTA

In []:

```
from transformers import RobertaTokenizer, RobertaForSequenceClassification
```

In []:

```
tokenizer = RobertaTokenizer.from_pretrained('roberta-base')
```

Sequence Length, 512

In []:

```
params['max_len'] = 512
params['batch_size'] = 8
```

In []:

```
train_dataloader, validation_dataloader = loaders(sentences_train, labels_train, sentences_test, labels_test, params['max_len'], params['batch_size'])
```

Truncation was not explicitly activated but `max_length` is provided a specific value, please use `truncation=True` to explicitly truncate examples to max length. Defaulting to 'longest_first' truncation strategy. If you encode pairs of sequences (GLUE-style) with the tokenizer you can select this strategy more precisely by providing a specific strategy to `truncation`.


```
/usr/local/lib/python3.6/dist-packages/transformers/tokenization_utils_base.py:2022:
FutureWarning: The `pad_to_max_length` argument is deprecated and will be removed in a future vers
ion, use `padding=True` or `padding='longest'` to pad to the longest sequence in the batch, or use
`padding='max_length'` to pad to a max length. In this case, you can give a specific length with `
max_length` (e.g. `max_length=45`) or leave max_length to None to pad to the maximal input size of
the model (e.g. 512 for Bert).
FutureWarning,
```

Original: I was wondering if anyone out there could enlighten me on this car I saw the other day. It was a 2-door sports car, looked to be from the late 60s/early 70s. It was called a Bricklin. The doors were really small. In addition, the front bumper was separate from the rest of the body. This is all I know. If anyone can tell me a model name, engine specs, years of production, where this car is made, history, or whatever info you have on this funky looking car, please e-mail.
Token IDs: torch.Size([512])

In []:

```
model = RobertaForSequenceClassification.from_pretrained(
    "roberta-base",
    num_labels = 20,
    output_attentions = False,
    output_hidden_states = False,
)
model.cuda();
```

Some weights of the model checkpoint at roberta-base were not used when initializing RobertaForSequenceClassification: ['lm_head.bias', 'lm_head.dense.weight', 'lm_head.dense.bias', 'lm_head.layer_norm.weight', 'lm_head.layer_norm.bias', 'lm_head.decoder.weight', 'roberta.pooler.dense.weight', 'roberta.pooler.dense.bias']

- This IS expected if you are initializing RobertaForSequenceClassification from the checkpoint of a model trained on another task or with another architecture (e.g. initializing a BertForSequenceClassification model from a BertForPreTraining model).
- This IS NOT expected if you are initializing RobertaForSequenceClassification from the checkpoint of a model that you expect to be exactly identical (initializing a BertForSequenceClassification model from a BertForSequenceClassification model).

Some weights of RobertaForSequenceClassification were not initialized from the model checkpoint at roberta-base and are newly initialized: ['classifier.dense.weight', 'classifier.dense.bias', 'classifier.out_proj.weight', 'classifier.out_proj.bias']

You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.

With Fine Tuning

In []:

```
params['epochs'], params['optimizer'], params['scheduler'] = model_parameters(model)
```

In []:

```
train_losses, valid_losses = train_loop(model, train_dataloader, validation_dataloader, params['opt
imizer'], params['scheduler'], params['epochs'], 1234, device)
```

```
Epoch 1/2, Train Loss: 1.2981    Valid Loss: 1.1378, Duration: 0:23:06.112439
Epoch 2/2, Train Loss: 0.7336    Valid Loss: 1.0544, Duration: 0:23:06.559929
```

In []:

```
train_acc = get_accuracy(train_dataloader, model, device)
valid_acc = get_accuracy(validation_dataloader, model, device)
print("Train accuracy:", train_acc, "Valid Accuracy:", valid_acc)
```

Train accuracy: 0.851864946084497 Valid Accuracy: 0.7130908125331917

Without Fine Tunning

In []:

```
model = RobertaForSequenceClassification.from_pretrained(
    "roberta-base",
    num_labels = 20,
    output_attentions = False,
    output_hidden_states = False,
)
model.cuda();
```

Some weights of the model checkpoint at roberta-base were not used when initializing RobertaForSequenceClassification: ['lm_head.bias', 'lm_head.dense.weight', 'lm_head.dense.bias', 'lm_head.layer_norm.weight', 'lm_head.layer_norm.bias', 'lm_head.decoder.weight', 'roberta.pooler.dense.weight', 'roberta.pooler.dense.bias']

- This IS expected if you are initializing RobertaForSequenceClassification from the checkpoint of a model trained on another task or with another architecture (e.g. initializing a BertForSequenceClassification model from a BertForPreTraining model).
- This IS NOT expected if you are initializing RobertaForSequenceClassification from the checkpoint of a model that you expect to be exactly identical (initializing a BertForSequenceClassification model from a BertForSequenceClassification model).

Some weights of RobertaForSequenceClassification were not initialized from the model checkpoint at roberta-base and are newly initialized: ['classifier.dense.weight', 'classifier.dense.bias', 'classifier.out_proj.weight', 'classifier.out_proj.bias']

You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.

In []:

```
for param in model.parameters():
    param.requires_grad = False
for param in model.classifier.parameters():
    param.requires_grad = True
```

In []:

```
params['epochs'], params['optimizer'], params['scheduler'] = model_parameters(model)
```

In []:

```
train_losses, valid_losses = train_loop(model, train_dataloader, validation_dataloader, params['optimizer'], params['scheduler'], params['epochs'], 1234, device)
```

```
Epoch 1/2, Train Loss: 2.9565    Valid Loss: 2.9206, Duration: 0:10:40.283758
Epoch 2/2, Train Loss: 2.8939    Valid Loss: 2.8862, Duration: 0:10:40.264726
```

In []:

```
train_acc = get_accuracy(train_dataloader, model, device)
valid_acc = get_accuracy(validation_dataloader, model, device)
print("Train accuracy:", train_acc, "Valid Accuracy:" ,valid_acc)
```

```
Train accuracy: 0.35310235106947147 Valid Accuracy: 0.3341741901221455
```

Conclusion:- After implementing all the experiments we observe that "ROBERTA" gives us the best possible validation accuracy for both fine-tuned model and without fine-tuned model of 33.4% & 71.3% respectively.

-----THE END-----